

Deep Learning - Generative Adversarial Text to Image Synthesis

Alex Valle (4854159) Gabriele Berruti (4791919)

Summary

The paper that has been chosen to be analysed is a 2016 Paper that showcases the effort of six researchers from the University of Michigan and Saarbrücken (Germany). The document shows how the use of GANs (Generative Adversarial Networks) allows for advancements in the generation of synthetic images starting from a textual description. It compares the method proposed by the research with previous architectures that, although far from the described goal, are capable of obtaining valid textual feature representations. In particular, the paper aims to demonstrate the effectiveness of the model in generating images of birds and flowers based on a precise textual description of them.

Introduction

In 2016, the ability of an AI system to generate realistic and coherent images from textual descriptions (such as "small red bird with a blue beak") was a current issue and far from being achieved. It should be noted that it's necessary to use natural language and domain-specific attributes to describe the image to be generated.

Datasets

Bird and Flower images from human-written descriptions .

- **Caltech-UCSD Birds Database** (Wah et al., 2011) (CUB) : Dataset used in related Papers previously described and in this one as well .
- **Oxford-102 Flowers Dataset** Dataset used in this paper that correspond to have 5 text descriptions per image .
- **Test Dataset - MS COCO Dataset** (Lin et al., 2014) In addition to birds and flowers more general images and text descriptions .

How to reach the goal ?

The difficulty of translating words into images may be divided into two subproblems.

1. : First, learn a feature vector from a specific text based on the visualization we want to obtain .
2. : Given these features through the use of a certain architecture, create a realistic and coherent image.

Effective Problems

The paper highlights an intrinsic problem: when trying to generate images from textual descriptions, there are many different configurations that can be correct. (Multi-modal Problem). And even the reverse task, that is, generating descriptions from images, is difficult, but it is easier because it can be handled by predicting one word at a time, based on what has

already been generated and the image itself.

Current Paper innovation ?

Our technique differs from previous conditional GANs by focusing on text descriptions rather than class labels. This is the first design capable of differentiating from character to pixel level. The paper develop a manifold interpolation regularizer for the GAN generator, which enhances sample quality, including "zero-shot" categories on CUB. It describes a model that generates 64x64 visually convincing pictures from text using a GAN. It differs from other models that just employ GANs for post-processing. In practice it used a character-level text encoder and class-conditional GAN and focus in implementing a new architecture and using it on fine-grained image datasets described before (CUB and Oxford Flowers) . Testing on MOCO dataset and test set disjoint from Training set can return a strong indicator on the performance of the system .

Background Knowledge

GAN

The GAN architecture consist of two principal component : Generator (G) and Discriminator (D) . The main goal of D is to distinguish between real training images and generated images coming from G . The idea is to maximize the logarithm of the loss of the Discriminator sampling images both from training set and generated while at the same time minimizing the outcoming of the log of 1 minus the loss of the Discriminator which has as input the generated image from random noise .

The Log function has a domain and an image that ranges from 0 to ∞ . It is possible to see that when the argument of the function tends to zero, the limit tends to $-\infty$, and when the argument tends to $+\infty$, the limit also tends to $+\infty$.




Figure 1: Log Function

Formula

$$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim p_{\text{data}}(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))] \quad (1)$$

- $D(x)$: the Discriminator as output layer use the Sigmoid function which map the data in a range from 0 to 1 .
- $G(x)$: the Generator as output layer return a dimension equal to \mathbb{R}^D

Therefore, based on these outputs and the behavior of the log function, it is possible to understand why the generator is minimized and the discriminator is maximized in the formula.

So in detail the log function of the first expected value must have as arguments elements between 0 and 1 at most (output of D), therefore the first part of the equation, trying to maximize D, sees mainly negative values and at most equal to 0.

At the same time, the second equation with the second expected value calculates the difference between $[1 - D(G(z))]$, and since the argument of D tends to $-\infty$ and we are simultaneously trying to maximize D and minimize G, the difference will result in a value between 0 and 1 (always positive).

Therefore, the limit of the logarithm as it approaches 0 will be equal to $-\infty$, while if the argument is 1, then the result is zero.

It's also possible to maximize G and use the $\log([D(G(z))])$ equation instead of $\log([1 - D(G(z))])$ and obtain the same result .

Text Encoder and Image Encoder

To obtain a vector representation of text descriptions the main paper used the approach of Reed in 2006 using Deep convolutional and symmetric architecture based on both images and its corresponding text description . The main idea is to have a dataset (Training set) composed of N elements :

$\{(v_n, t_n, y_n) : n = 1, \dots, N\}$
in which :

- v_n : correspond to the N images
- t_n : correspond to the N text description
- y_n : correspond to the class label (usually there are more than two usually M)

First of all the paper define the classifier for both images and text description :

$$f_v(v) = \arg \max_{y \in \mathcal{Y}} \mathbb{E}_{t \sim \mathcal{T}(y)} [\phi(v)^T \varphi(t)] \quad (2)$$

$$f_t(t) = \arg \max_{y \in \mathcal{Y}} \mathbb{E}_{v \sim \mathcal{V}(y)} [\phi(v)^T \varphi(t)] \quad (3)$$

So in the first equation from all the possible class labels (y) we are trying to find the one that maximize the correlation between the input image and all the text descriptions on all the classes .

To do it we compute the multiplication between the encoded version of the input image v and the current text description (t) related to the current class .

In practice $\phi(v)$ and $\varphi(t)$ return a vector and we compute the similarity via a vector multiplication that return a scalar .

To compute the expected value in practice we compute the mean of all the scalar that we obtain from this multiplication on the specific class (y) .

In the second equation the formula is similar but on

the input text description (t) . The main idea was that an encoded description should have a higher compatibility with the images related to its class and so obtain higher value

- \mathcal{Y} : Set of the classes in which the images are divided .
- $\mathcal{T}(y)$: Set of the text description related to the y class .
- $\mathcal{V}(y)$: Set of the images related to the y class .
- $f_v(v)$: Classifier based on the images which return the class with higher compatibility with the image v
- $f_t(t)$: Classifier based on text description which return the class with higher compatibility with the description t
- $\phi(v)$: Image Encoder representation on v already trained on different dataset which return a vector (es. Obtained via CNN) .
- $\varphi(t)$: Text Encoder representation on t already trained on different dataset which return a vector (es. Obtained via CNN or LSTM).
- Δ : Hinge or 0-1 loss for classification defined in which if the argument is higher than 0 it return 0 or 1 otherwise
- $\mathbb{E}_{t \sim \mathcal{T}(y)}$: Expected value in which t correspond to all the text description related at the y class .
- $\mathbb{E}_{v \sim \mathcal{V}(y)}$: Expected value in which v correspond to all the images related at the y class .

Hinge or 0-1 Loss in this case is:

$$\Delta(y, f(x)) = \max(0, 1 - y * f(x)) \quad (4)$$

$$\varphi_{loss} = \frac{1}{N} \sum_{n=1}^N \Delta(y_n, f_v(v_n)) + \Delta(y_n, f_t(t_n)) \quad (5)$$

The main idea to compute the φ_{loss} loss function was to learn the weights of the Text Encoder (φ) and fine-tuning it on our dataset .

Our implementation of Text/Image Encoding via CLIP

During our implementation part instead than using the Text Encoder used by the Paper we decided to fine-tune the CLIP (Contrastive Language-Image Pre-Training , developed by OpenAI) Text Encoder on the complete CUB and Flower dataset for few epochs, in such a way to obtain relevant vector representation .

Methodology

All the methods are based on a GAN architecture with a Generator G and a Discriminator D . And use (φ) as Text Encoder that is a character-level convolutional- recurrent neural network .

The generator is defined as :

$$G : \mathbb{R}^Z \times \mathbb{R}^T \rightarrow \mathbb{R}^D,$$

And the discriminator is defined as :

$$D : \mathbb{R}^D \times \mathbb{R}^T \rightarrow \{0,1\},$$

where

- D : is the dimension of the generated image
- Z : is the dimension of the input noise
- T : is the dimension of the text emdedded with φ

- Generator Part -

First of all the generated image is defined as :

$$\hat{x} \leftarrow G(z, \varphi(t))$$

This Generator take as input a noise vector of dimension Z (In our implementation $Z = 100$) whose values are sampled from a Gaussian distribution between 0 and 1

$$z \in \mathbb{R}^Z \sim \mathcal{N}(0,1)$$

and the vector coming from the Text Encoder φ for the specific text description t obtaining $\varphi(t)$.

First part - Projection

The first step of the Generator Net is to reduce the dimension of $\varphi(t)$ using one fully-connected layer to project to the dimension of the projection P .

This part continue using a Normalization Layer (BatchNorm1d) and a Leaky-Relu as activation function with a specific negative slope . So the values on the vector can be negative and positive depending on the situation .

In our implementation φ correspond to CLIP while T is equal to 512 and P is equal to 128

After this first part an 128-dim vector is obtained and defined as :

$$p = projection(\varphi(t))$$

Second part - Concatanation

To compose the Latent Space vector (LSV) the paper suggested to concatenate the noise vector z with the projected vector p .

In our specific implementation the latent space has a dimension of H ($H=228$) which is simple the concatenation of the two vector of dimension $T=100$ and $P=128$.

$$h = concat(z, p)$$

Third part - Final Net to generated the image 64x64 RGB As input of the Final Net part of the Generator is the Latent Space Vector (LSV) of dimension H ($H=228$).

This Final Net is composed of 4 set of layers in which the first one is a Convolutional Transpose 2D Layers ("to up-sample") followed by a Normalization layer and Relu as activation function so we only have values greater than 0 .

In our implementation the initial dimension of the vector was (228,1,1) but after only the first Convolutional Transpose layer it is projected into an higher dimension of (512,8,8) .

After using the other 3 set of layers the dimension will be equal to (64,32,32) . The last layer is useful because permit to obtain an image of the shape (3,64,64) and

is composed of a Convolutional Transpose 2D Layer and Tanh as activation function .

Using Tanh with in input values coming from the last layer that can be positive and negative as well due to the weight of the 2D Convolutional Transpose Layer we will obtain an image 64x64 with 3 channel RGB with values that ranges from -1 to 1 due to Tanh .

$$\hat{x} = generate(h)$$

And G is defined as this 3 Parts applied in order one after another .

- Discriminator Part -

The Discriminator is defined as this :

$$D(\hat{x}, \varphi(t)) = \{x_i \in [0, 1] \mid i = 0, 1, \dots, 120\}$$

It takes in input the generated image \hat{x} and depending on the situation an encoded version of the text description $\varphi(t)$ related or not to image generated .

Similarly to what we have seen in the Generator Part the Discriminator is sub-divided in 3 Parts or Networks .

First Net - Down-Sampling The First Part is composed of a several (in our implementation 4) set of Convolutional 2D Layers ("to down-sample") followed by a normalization Layer (BatchNorm2d except the first set) and Leaky-Relu as activation function . After passing through this Net the generated image of a shape of (3,224,244) will have a dimension of (512,14,14) .

$$q = DownSampling(\hat{x})$$

Second Part - Projection Text Embedding and Concatenation

This part consist of two phases and it take as argument the down-sampled generated image q and the encoded version of the text description $\varphi(t)$. The first phase consist in projecting $\varphi(t)$ in a lower-dimension P (P=128) using one fully-connected layer and normalization layer followed by a Leaky-Relu as activation

function in the exact same way that we have done in the Generator . After this first phase a 128-dim vector is obtained and defined as :

$$p = projection(\varphi(t))$$

The last task of the first phase was to adapt the project dimension of p (that is (128,1,1)) in such a way to be compatible with the second phase which concatenate p with q . So the dimension of p is squeezed into a dimension that allow to concatenate the two data and it became (128,14,14)

$$\hat{p} = squeeze(p)$$

The second Phase consist of the concatenation of q and \hat{p} which result in a dimension of (512+128, 14, 14) = (256, 640, 14, 14) This concatenation permit to obtain the Latent vector c :

$$c = concatenate(q, \hat{p})$$

Third Part - Final Net This Final Net it's only composed of One Convolutional 2D Layer and Sigmoid as activation function which map the input in a range that goes from 0 to 1 . The first Convolution Layer modify the shape from (256, 640, 14, 14) to (512,4,4) while the sigmoid function change its shape into (1,11,11) . At the end of this phase we flatten the 11x11 values into a vector of 121-dimension . So at the end for each generated image and text Embedding we obtain a 121-dim vector representing the discriminator factors .

$$d_{loss} = FinalNet(c)$$

Architectures

(Vanilla GAN) : Generative Adversarial Network without Conditional Latent Space

The Vanilla implementation stands for an implementation in which the Generator only depend on the noise

vector z and not depend on the Conditional Latent Space $\varphi(t)$ and the same stand for the discriminator which depends only on the actual or the generated image .

(GAN-CLS) : Generative Adversarial Network with Conditional Latent Space

The idea behind this architecture is to use φ as Text Encoder (CLIP) to create the embeddings related to text description t coherent with the image (x) to obtain h and one not correlated to the image sampled randomly \hat{t} to obtain \hat{h} . Another step is to generate the Noise Vector z randomly sampled from a Gaussian Distribution .

The following step is to generate an image passing through z and h as explained in the Generator Part and obtaining \hat{x} . As well the Discriminator have to compute a value s_r that correspond to the loss passing through it the real image z and the corresponding text emdedding h .

In the same way we obtain s_r passing at inference time the image x and \hat{t} and we calculate s_f passing the image x and h . The Final part consist in computing the Discriminator loss L_D as the $\log(s_r) + (\log(1 - s_w) + \log(1 - s_f)) / 2$.

In the actual implementation the loss of the Discriminator is computed using the BCELoss (Binary Cross Entropy Loss) between s_r and a "smoothed" version of the label filled with 1 computing " d_loss_r " and again two times , first with s_f and a "fake label" filled with 0 computing " d_loss_f " and after using s_f and the "fake label" computing " d_loss_f " .

So the final value is computed summing the 3 obtained values :

$$L_D = d_loss_f + d_loss_w + d_loss_r$$

After this computation the Discriminator is updated using backpropagation based on the gradient with α as parameter .

For the Generator Loss L_D we don't compute only the logarithm of s_f (in Algorithm 1) but we use a custom loss composed of 3 part . The first one compute the BCELoss between s_f and the "real label" which is a vector filled with 1 .

The second one compute the MSELoss (Mean Squared Error Loss) between q_r (after Down-Sampling) and q_f coming from the Discriminator of s_r and s_f . The Third part is computed the L1Loss between the generated image \hat{x} and the real image x . The 3 part are summed to obtain L_G

$$L_G = BCELoss(s_f, real_label) + MSELoss(q_r, q_f) + L1Loss(\hat{x}, x)$$

And finally after this computation the Generator is updated using backpropagation based on the gradient with α as parameter .

Algorithm 1 GAN-CLS training algorithm with step size α , using minibatch SGD for simplicity.

Require: minibatch images x , minibatch matching text t , minibatch mis-matching text \hat{t} , number of training batches S

for $n = 1$ to S **do**

$h \leftarrow \varphi(t)$ {Encode matching description}

$\hat{h} \leftarrow \varphi(\hat{t})$ {Encode mis-matching description}

$z \sim \mathcal{N}(0, 1)^Z$ {Extract Noise Vector}

$\hat{x} \leftarrow G(z, h)$ {Forward through generator}

$s_r \leftarrow D(x, h)$ {Real image, right text}

$s_w \leftarrow D(x, \hat{h})$ {Real image, wrong text}

$s_f \leftarrow D(\hat{x}, h)$ {Fake image, right text}

$L_D \leftarrow \log(s_r) + (\log(1 - s_w) + \log(1 - s_f)) / 2$

$D \leftarrow D - \alpha \nabla L_D$ {Update discriminator}

$L_G \leftarrow \log(s_f)$

$G \leftarrow G - \alpha \nabla L_G$ {Update generator}

end for

(GAN-INT) : Generative Adversarial Network Interpolated

The main difference with respect with the previous implementation relies in the fact that Generation of the "fake image" do not depend not anymore only on z and t , but in an interpolation of two different text emdeddings . In practice the paper found that fixing $\beta = 0.5$ works well and want to underline that t_1 and t_2 may come from different images and even different categories. This parte has not been implemented and

tasted but with little modification on the actual code can be performed .

$$\mathbb{E}_{h1, h2 \sim p_{data}} [\log(1 - D(G(z, \beta h1 + (1 - \beta)h2)))] \quad (6)$$

- WGAN : - Wessertstain GAN

This typology of GAN is not present in the Paper but follow the same concept described before . In principle we can observe first all the algorithm :

Algorithm 2 WGAN training algorithm with step size α , using minibatch SGD for simplicity.

Require: minibatch images x , minibatch matching text t , minibatch mis-matching text \tilde{t} , number of training batches S , number of iteration for the Discriminator NI

```

for  $n = 1$  to  $S$  do
   $h \leftarrow \varphi(t)$            {Encode matching description}
  for  $n = 1$  to  $NI$  do
     $z \sim \mathcal{N}(0, 1)^Z$        {Extract Noise Vector}
     $\hat{x} \leftarrow G(z, h)$     {Forward through generator}
     $s_r \leftarrow D(x, h)$     {Real image, right text}
     $s_f \leftarrow D(\hat{x}, h)$   {Fake image, right text}
     $L_D \leftarrow (s_r) - (s_f)$ 
    Clip weights of  $D$  within  $[-c, c]$ 
    {Lipschitz constraint}
  end for
   $z \sim \mathcal{N}(0, 1)^Z$        {Extract again Noise Vector}
   $\hat{x} \leftarrow G(z, h)$     {Forward through generator}
   $s_f \leftarrow D(\hat{x}, h)$   {Fake image, right text}
   $L_G \leftarrow -(s_f)$  {Wasserstein loss for Generator}
end for

```

The idea behind this architecture is to use φ as a Text Encoder to create embeddings related to a text description t that matches the image x , obtaining h . This embedding h serves as the condition for both the Generator and Discriminator.

Another step involves sampling a Noise Vector z randomly from a Gaussian Distribution $\mathcal{N}(0, 1)^Z$.

The following step is to generate a fake image \hat{x} by passing z and h through the Generator G .

The Discriminator D computes a score s_r by passing

the real image x and the corresponding text embedding h , which reflects how well D recognizes the real data.

Similarly, the Discriminator computes a score s_f by passing the generated image \hat{x} and the same text embedding h .

The Discriminator loss L_D is calculated as the difference $s_r - s_f$, which represents the Wasserstein distance. To enforce the Lipschitz constraint, the weights of the Discriminator are clipped within a fixed range $[-c, c]$ after every update.

This ensures the Discriminator satisfies the required gradient properties for stable training. After training the Discriminator for multiple steps, the Generator is trained.

A new Noise Vector z is sampled from $\mathcal{N}(0, 1)^Z$, and a fake image \hat{x} is generated by passing z and h through G .

The Discriminator then computes a new score s_f for this generated image, and the Generator loss L_G is calculated as $-s_f$.

This encourages the Generator to produce images that maximize the Discriminator's score, effectively improving the quality of the generated images.

The training process alternates between optimizing the Discriminator and the Generator. Over multiple iterations, the Discriminator learns to distinguish real and fake images better, while the Generator learns to produce more realistic images.

This iterative process ensures that the Wasserstein distance between the real and generated data distributions is minimized, leading to stable and efficient GAN training.

Generator Inversion for Style Transfer

The text encoding $\varphi(t)$ represents the content of an image (e.g., flower shape and colors). To generate realistic images, the noise sample z should encode style elements such as background color and pose. Using a trained GAN, it becomes possible to transfer the style from a query image to align with the content of a specific text description.

This is achieved by training a convolutional network

S to invert the generator G , allowing the recovery of z from generated samples $\hat{x} \leftarrow G(z, \varphi(t))$. The style encoder S is trained using a simple squared loss function:

$$\mathcal{L}_{\text{style}} = \mathbb{E}_{t, z \sim \mathcal{N}(0,1)} \|z - S(G(z, \varphi(t)))\|_2^2 \quad (6)$$

Here, S denotes the style encoder network.

Once the generator and style encoder are trained, the style from a query image x can be transferred to match a text description t as follows:

$$s \leftarrow S(x), \quad \hat{x} \leftarrow G(s, \varphi(t))$$

In this process, s represents the extracted style, and \hat{x} is the final image output.

Experiment conducted by the Paper

Datasets

CUB Dataset

The CUB dataset contains 11,788 bird images belonging to 200 distinct categories.

Oxford-102 Dataset

The Oxford-102 dataset consists of 8,189 flower images, grouped into 102 categories.

Data Splits

For training and evaluation, datasets are split into class-disjoint subsets:

- **CUB:** 150 categories are used for training and validation, while the remaining 50 are used for testing.
- **Oxford-102:** 82 categories are assigned to training and validation, and 20 categories are reserved for testing.

Captions

Each image in both datasets is accompanied by 5 captions. During training, a random image view (e.g., crop, flip) and one randomly selected caption are used for each mini-batch.

Text Encoder

Encoder Architecture

The text encoder employs a deep convolutional-recurrent network, combining a character-level ConvNet with a recurrent neural network (char-CNN-RNN). This architecture generates 1,024-dimensional embeddings for textual descriptions.

Text Embedding

Text captions are embedded into a 1,024-dimensional space via structured joint embedding with GoogLeNet features.

Pre-training

Pre-training the text encoder accelerates the training of the generator and discriminator, enabling faster experimentation. While pre-training is not a strict requirement, end-to-end training results are provided in the supplement for completeness.

Generalization

Qualitative results from the MS COCO validation set illustrate the approach’s ability to generalize beyond the datasets used in training.

GAN Architecture

Training Image Size

All training images are resized to $64 \times 64 \times 3$.

Text Feature Projection

Text embeddings are projected to a 128-dimensional space before being concatenated with convolutional feature maps in both the generator and discriminator networks.

Training Process

The generator and discriminator in the GAN-CLS architecture are updated in alternating steps to optimize the GAN architecture.

Hyperparameters

The following hyperparameters are used during training:

- Base learning rate: 0.0002.
- Optimizer: ADAM, with a momentum parameter of 0.5.
- Generator noise: Sampled from a 100-dimensional unit normal distribution.
- Mini-batch size: 64.
- Number of epochs: 600.

Implementation

The implementation of the model is based on the `drgan.torch2` framework.

Paper Qualitative Results

Comparison of GAN Variants

We compare the following GAN architectures:

- **GAN Baseline:** The basic GAN model without any specific improvements for text-image matching.
- **GAN-CLS:** A GAN model incorporating an image-text matching discriminator.
- **GAN-INT:** A GAN variant utilizing text manifold interpolation.
- **GAN-INT-CLS:** A model combining both text manifold interpolation and the image-text matching discriminator.

Results on the CUB Dataset

Qualitative results for the CUB dataset:

- The **GAN Baseline** and **GAN-CLS** models correctly reproduce some color information. However, the generated images do not appear realistic.
- **GAN-INT** and **GAN-INT-CLS** models produce plausible bird images that match either all or part of the captions.
- Additional robustness analysis for each GAN variant on the CUB dataset is provided in the supplement.

Results on the Oxford-102 Dataset

Qualitative results for the Oxford-102 Flowers dataset are presented in Figure 4:

- All four models are capable of generating plausible flower images that align with their respective captions.
- The **GAN Baseline** exhibits the highest variety in flower morphology, especially when the caption does not specify petal types.
- Other models, such as **GAN-CLS**, **GAN-INT**, and **GAN-INT-CLS**, generate more class-consistent flower images.
- It is speculated that generating flowers is easier than birds due to structural regularities in bird species, making it simpler for the discriminator to identify fake birds compared to fake flowers.

Additional Results

Supplementary materials include additional examples for the following:

- **GAN-INT** and **GAN-INT-CLS** models on both CUB and Oxford-102 datasets.
- **Vanilla GAN:** An end-to-end variant of GAN-INT-CLS that does not rely on pre-training the text encoder $\phi(t)$.

Disentangling Style and Content

This section explores the capability of the model to separate style from content.

By *content*, we refer to the intrinsic visual characteristics of the bird, such as the shape, size, and color of its body parts. *Style*, on the other hand, encompasses external factors like background color and pose orientation.

Since the text embedding primarily encodes content information and usually excludes style details (e.g., captions rarely describe background or pose), the GAN must utilize the noise vector z to account for variations in style.

To generate realistic images, the disentanglement of these factors is crucial.

To measure the extent of style-content disentanglement on the CUB dataset, we defined two tasks:

pose verification and background color verification. Each task required constructing pairs of similar and dissimilar images. Style vectors for these images were obtained by passing them through a style encoder network, trained to invert the generator’s outputs back to the noise vector z .

If style and content are disentangled, images with similar styles (e.g., same pose) should have higher similarity scores than images with different styles (e.g., different poses).

To recover z , we inverted the generator networks following the procedure in Subsection 4.4. Verification pairs were created by clustering images into 100 groups using K-means. For background color verification, clustering was performed based on the average RGB values of the background.

For pose verification, clustering relied on six keypoint coordinates (beak, belly, breast, crown, forehead, and tail).

Evaluation was performed by calculating predicted style vectors for image pairs using style encoders for GAN, GAN-CLS, GAN-INT, and GAN-INT-CLS models.

Similarity scores were computed with cosine similarity, and the AU-ROC metric was reported, averaged over five folds.

As a baseline, cosine similarity between text features from the text encoder was also calculated.

The results, shown in Figure 5, confirm that captions alone do not provide style-related information. Consistent with qualitative observations, models incorporating interpolation regularization (GAN-INT and GAN-INT-CLS) achieved superior performance for these tasks.

Specifically:

- **Pose Verification:** ROC curves demonstrate that style encoders effectively distinguish between similar and different poses.
- **Background Color Verification:** ROC curves illustrate that models can separate images based on background color variations.

Pose and Background Style Transfer

GAN-INT-CLS, when combined with a trained style encoder (as described in the Style Transfer subsection), enables style transfer from an unseen query image to a new text description.

Remarkably, this approach often retains intricate background details, such as a tree branch on which the bird is perched, showcasing the effectiveness of disentangling style and content.

The ability of GAN-INT-CLS to disentangle style is particularly noteworthy as it facilitates a straightforward method of generalization.

By leveraging this disentanglement, the model can combine previously observed content (e.g., text descriptions) with previously seen styles in novel pairings, thereby generating realistic and plausible images that are entirely distinct from those seen during training.

Another pathway for generalization involves utilizing attributes present in the training data (e.g., "blue wings" and "yellow belly") to create new combinations.

Sentence Interpolation

Although ground-truth text descriptions are not available for the interpolated points, the generated images

appear visually coherent and plausible.

By holding the noise distribution constant, the only variable factor across each row is the text embedding. This approach allows interpolations to accurately capture content-related changes, such as a bird transitioning from blue to red, while preserving consistent pose and background elements.

In addition to sentence interpolation, Figure 8 (Right) presents results achieved through noise interpolation. In this case, two random noise vectors are sampled, and by keeping the text embedding fixed, we interpolate between these noise vectors.

The resulting images display a smooth style transition while maintaining fixed content.

This demonstrates the model’s ability to disentangle style and content effectively, enabling independent manipulation of each aspect to produce seamless variations in generated images.

Training on other dataset - MOCO

To demonstrate the generalization capability of our approach, we trained a GAN-CLS model on the MS-COCO dataset.

We have to underline that unlike CUB and Oxford-102, MS-COCO features diverse images containing multiple objects and variable backgrounds.

Despite this complexity, we employed the same text encoder, GAN architecture, and hyperparameters (learning rate, mini-batch size, and number of epochs) used in the previous datasets.

The main difference lies in the text encoder training, as COCO lacks a single object category per class, requiring an instance-level image and text matching approach.

Figure 7 displays examples of generated images alongside their corresponding ground-truth captions.

The results exhibit sharpness typical of GAN-based synthesis methods and notable diversity in samples, achieved by varying the noise vector while keeping the text embedding fixed.

However, closer inspection reveals some limitations in scene coherence, particularly in complex scenarios like human figures in baseball scenes, where articulated parts are missing.

Future work could explore incorporating hierarchical

structures into the synthesis model to better handle multi-object scenes.

Additionally, a qualitative comparison with AlignDRAW (Mansimov et al., 2016) is provided in the supplement.

While GAN-CLS produces sharper and higher-resolution outputs roughly matching the query, AlignDRAW better captures fine-grained single-word variations. Extending the GAN-CLS generator network with temporal structures could improve its ability to handle nuanced text differences.

Paper Conclusions

This work presented a simple yet effective model for generating images from detailed visual descriptions. Our approach successfully synthesized multiple plausible visual interpretations for a given text caption. The introduction of a manifold interpolation regularizer significantly enhanced text-to-image synthesis on the CUB dataset.

We also demonstrated the ability to disentangle style and content, enabling bird pose and background transfer from query images to text descriptions.

Furthermore, our model showed generalizability to more complex scenarios, generating images with multiple objects and variable backgrounds on the MS-COCO dataset.

Future work will focus on scaling the model to produce higher-resolution images and incorporating a broader range of text inputs.

Code implementation

The code has been written in Python, while back in 2016 the researcher decided to write the code for training the Net in LUA. So first of all we understand the code written in LUA by the authors and after we decide to translate it in Python. The main reason why we do it is that the code was deprecated after about 10 years (especially the library) and for us was not possible to try to train the Net ourself. Talking about the Datasets they were also not reachable because the hyperlink redirect to a website

not existing anymore . So we try to develop a solution in which we train the Net with the same concept described before but using different Datasets . Also the trained weights cannot be adapted to the new solution that we develop , so we decide so re-train the Network from zero for about 200 epochs . After the 200 epochs we decide to stop the training because it was taking too long .

The code is subdivided in different part and diffent .ipynb file . The first one is "1) dataset usage.ipynb" in which is possible the are two example on how to load Bird and Flowers Dataset using the module "gan_t2i" and store it in HDF5 format other than using a Dataloader on this datasets . The second file is "2) CLIP - Fine Tuning.ipynb" in which is shown how to load the CLIP model (ViT-B/32) and to train it (we use an already trained CLIP network to extract text features) The Third file is "3.2) COLAB GAN example.ipynb" . In this file we combine all the module that we develop until now . First of all we decide which network we would like to train between : "Vanilla GAN", "GAN_INIT_CLS" and "WGAN" . After this part we download the weights related to the CLIP Network used to extract the text features and also the model itself . What about the Dataset initially it is stored in HDF5 format but this time is also transformed and normalized other than tokenized . After this section we create the training, validation and test dataloaders and check the outcome size of the CLIP model feature related to text and images . And then depending on the Newtwork that we would like to train we define it and define the embedding projection dimension that in our case is 128 . Finally after this long pre-processing and initialization part we can decide if start to train the Net using the correspondent Algoritm from zero or from a specific checkpoint (which corrispond to the weights) loading it in the Model . The result on the GAN model after 186 epochs of training from scratch are not enough to prodouce a result which correspond to the description given by the text , that's due to the fact that the Net is too complex and in fact to achieve a satisfactory result, at least 600 epochs are necessary, as the original paper also specifies. Besides the visualization factor it's possible to observe that the loss related to the Discrimination and Generator

part tend to move in the right direction , in fact the Generator loss especially in the WGAN model tend do decrease (minimize) and the Discriminator loss it's maximizing it's values as well as described in the Min-Max optimization formula .

WGAN on FASHION-MNIST Dataset

We try also to conduct some experiments on training the Net addressing the problem to other Datasets such as Fashion MNIST dataset (`keras.datasets.fashion_mnist`) .

So in general the WGAN employs the Wasserstein distance to define a value function with properties theoretically superior compared to the one introduced in the original GAN paper . As seen before to verify that the Discriminator (critic) respects the 1-Lipschitz requirement, the authors implemented weight clipping. However, this technique might lead to problems like as low convergence in deep critics and other unwanted phenomena. WGAN-GP substitutes weight clipping with a "gradient penalty." This variant includes a loss term to keep the L2 norm of the discriminator gradients around a set value, allowing for smoother training . But it's based on the Algorithm 2 flow-chart described in the section related to the Architectures.

Data Preparation

The dataset used is the FASHION-MNIST dataset , in which each sample is a 28x28 grayscale pixel image with values normalized in the range $[-1,1]$. The label related to the image will not be used in this case to train the WGAN .

Discriminator Net

In this case the Discriminator use a Zero Padding layer to transform the input image (28,28,1) into a shape of (32,32,1) and other 4 Convolutional Block to transform it in (16,16,64) \rightarrow (8, 8, 128) \rightarrow (4, 4, 256) \rightarrow (2, 2, 512) Leaky Relu is used as activation function .

The Convolutional Block is composed of 2D Convolutional Kernel and depending on the parameters a Normalization and dropout layer . The output shape of the discriminator is one value that represent the

"real" or "fake" classification of the input image, indicating whether the image is from the real dataset or generated by the Generator.

Generator Net

The Genetor Net use the `upsample_block(..)` function to handle the upsampling process by increasing the spatial dimensions of the input with `UpSampling2D`, followed by a `Conv2D` layer . Depending on the parameter it may also include `BatchNormalization` and a specific activation function (like `LeakyReLU` or `Tanh`) to introduce non-linearity and improve training stability other than a `Dropout` layer to prevent overfitting. The Generator starts by taking a noise vector as input, transforming it through a `Dense` layer into a tensor of shape (4, 4, 256). The data is then passed through three Upsampling blocks . In the same way in the second block, the shape changes from (8, 8, 128) to (16, 16, 64), again using `UpSampling2D` and `Conv2D`, with `BatchNormalization` and `LeakyReLU` applied. The third one increases the shape from (16, 16, 64) to (32, 32, 1) by applying `UpSampling2D` and `Conv2D` to reduce the channels to 1, followed by a `Tanh` activation to ensure the output is in the correct range. Lastly , a `Cropping2D` layer is applied to reduce the spatial dimensions from (32, 32) to (28, 28) to generate the image .

WGAN-GP MODEL

The model overrides the `keras.Model` module and overrides the `train_step` function derived from the model. The model consists of a discriminator and a generator. The discriminator processes real and fake images, while the generator creates fake images from random noise. The model is initialized with specific hyperparameters like `latent_dim` (dimension of the random input to the generator), `discriminator_extra_steps` (extra steps to train the discriminator), and `gp_weight` (weight for the gradient penalty).

Gradient Penalty Function

The model uses a custom gradient penalty function, which helps to enforce the Lipschitz constraint by calculating the norm of the gradient of the discriminator's output with respect to interpolated images. This is added to the discriminator's loss.

Training Step

During training, we work using batches of 512. The discriminator is updated multiple times per generator step (3 extra steps are typically used). For each discriminator step, fake images are generated based on the noise vector. The discriminator is used two times on the fake and real images, and the discriminator's loss is computed based on the function passed (we will define it after), obtaining `d.cost`. Based on the real and fake images, we compute the gradient penalty. Finally, the total discriminator loss (`d.loss`) is computed as :

```
d.loss = d.cost + gp * self.gp.weight.
```

The generator is trained after the discriminator. The generator's loss is computed based on the discriminator's output, and it depends on the generated images coming from the Generator Network. The generator loss function will be defined after. Each model (discriminator and generator) is updated using their respective optimizers. The training process alternates between updating the discriminator and the generator, with the gradient penalty ensuring stability. The optimizers used for this task are the same and correspond to Adam with a learning rate of 0.0002.

Discriminator Loss

The discriminator loss is simple and computes the mean of the batch values obtained on real and fake images, as well as the difference between them.

Generator Loss

The generator loss is straightforward because it is the negative mean of the logits (unnormalized output values produced by the model) for fake images,

pushing the generator to produce images that the discriminator classifies as real. A positive score suggests the image is real, and a negative score suggests it is fake. By minimizing this negative loss, the generator is maximizing the discriminator's score for the generated images, effectively making the fake images more "realistic" in the eyes of the discriminator.

Training time

The network is trained for around 30 epochs, with a noise dimension of 128. After each epoch, a callback generates and saves a number of images based on the latent noise vector as PNG files. The values from the generator are in the range of $[-1, 1]$ due to the `Tanh` activation function, so these values are rescaled to the range $[0, 255]$ to be saved correctly. We perform the inverse process to normalize the pixels in the range $[-1, 1]$ when preparing each sample of the dataset.




Figure 2: Discriminator Loss during Training (WGAN)

In general, especially because the problem has been simplified with respect to the previous one, after only 50 epochs we can appreciate good result in term of reconstruction of image related to the MNIST datasets. And thanks to the gradient penalty term we can generalize the enough to obtain a good reconstruction also on "zero-shot" data coming from the same dataset.

The visual results of the model's performance after various epochs are shown below. As it is possible to observe, the reconstructions generated by the "Gen-

erator” become increasingly coherent with the noise given as input as the epochs increase.

Figure 3: *Grid of generated images at different epochs*

